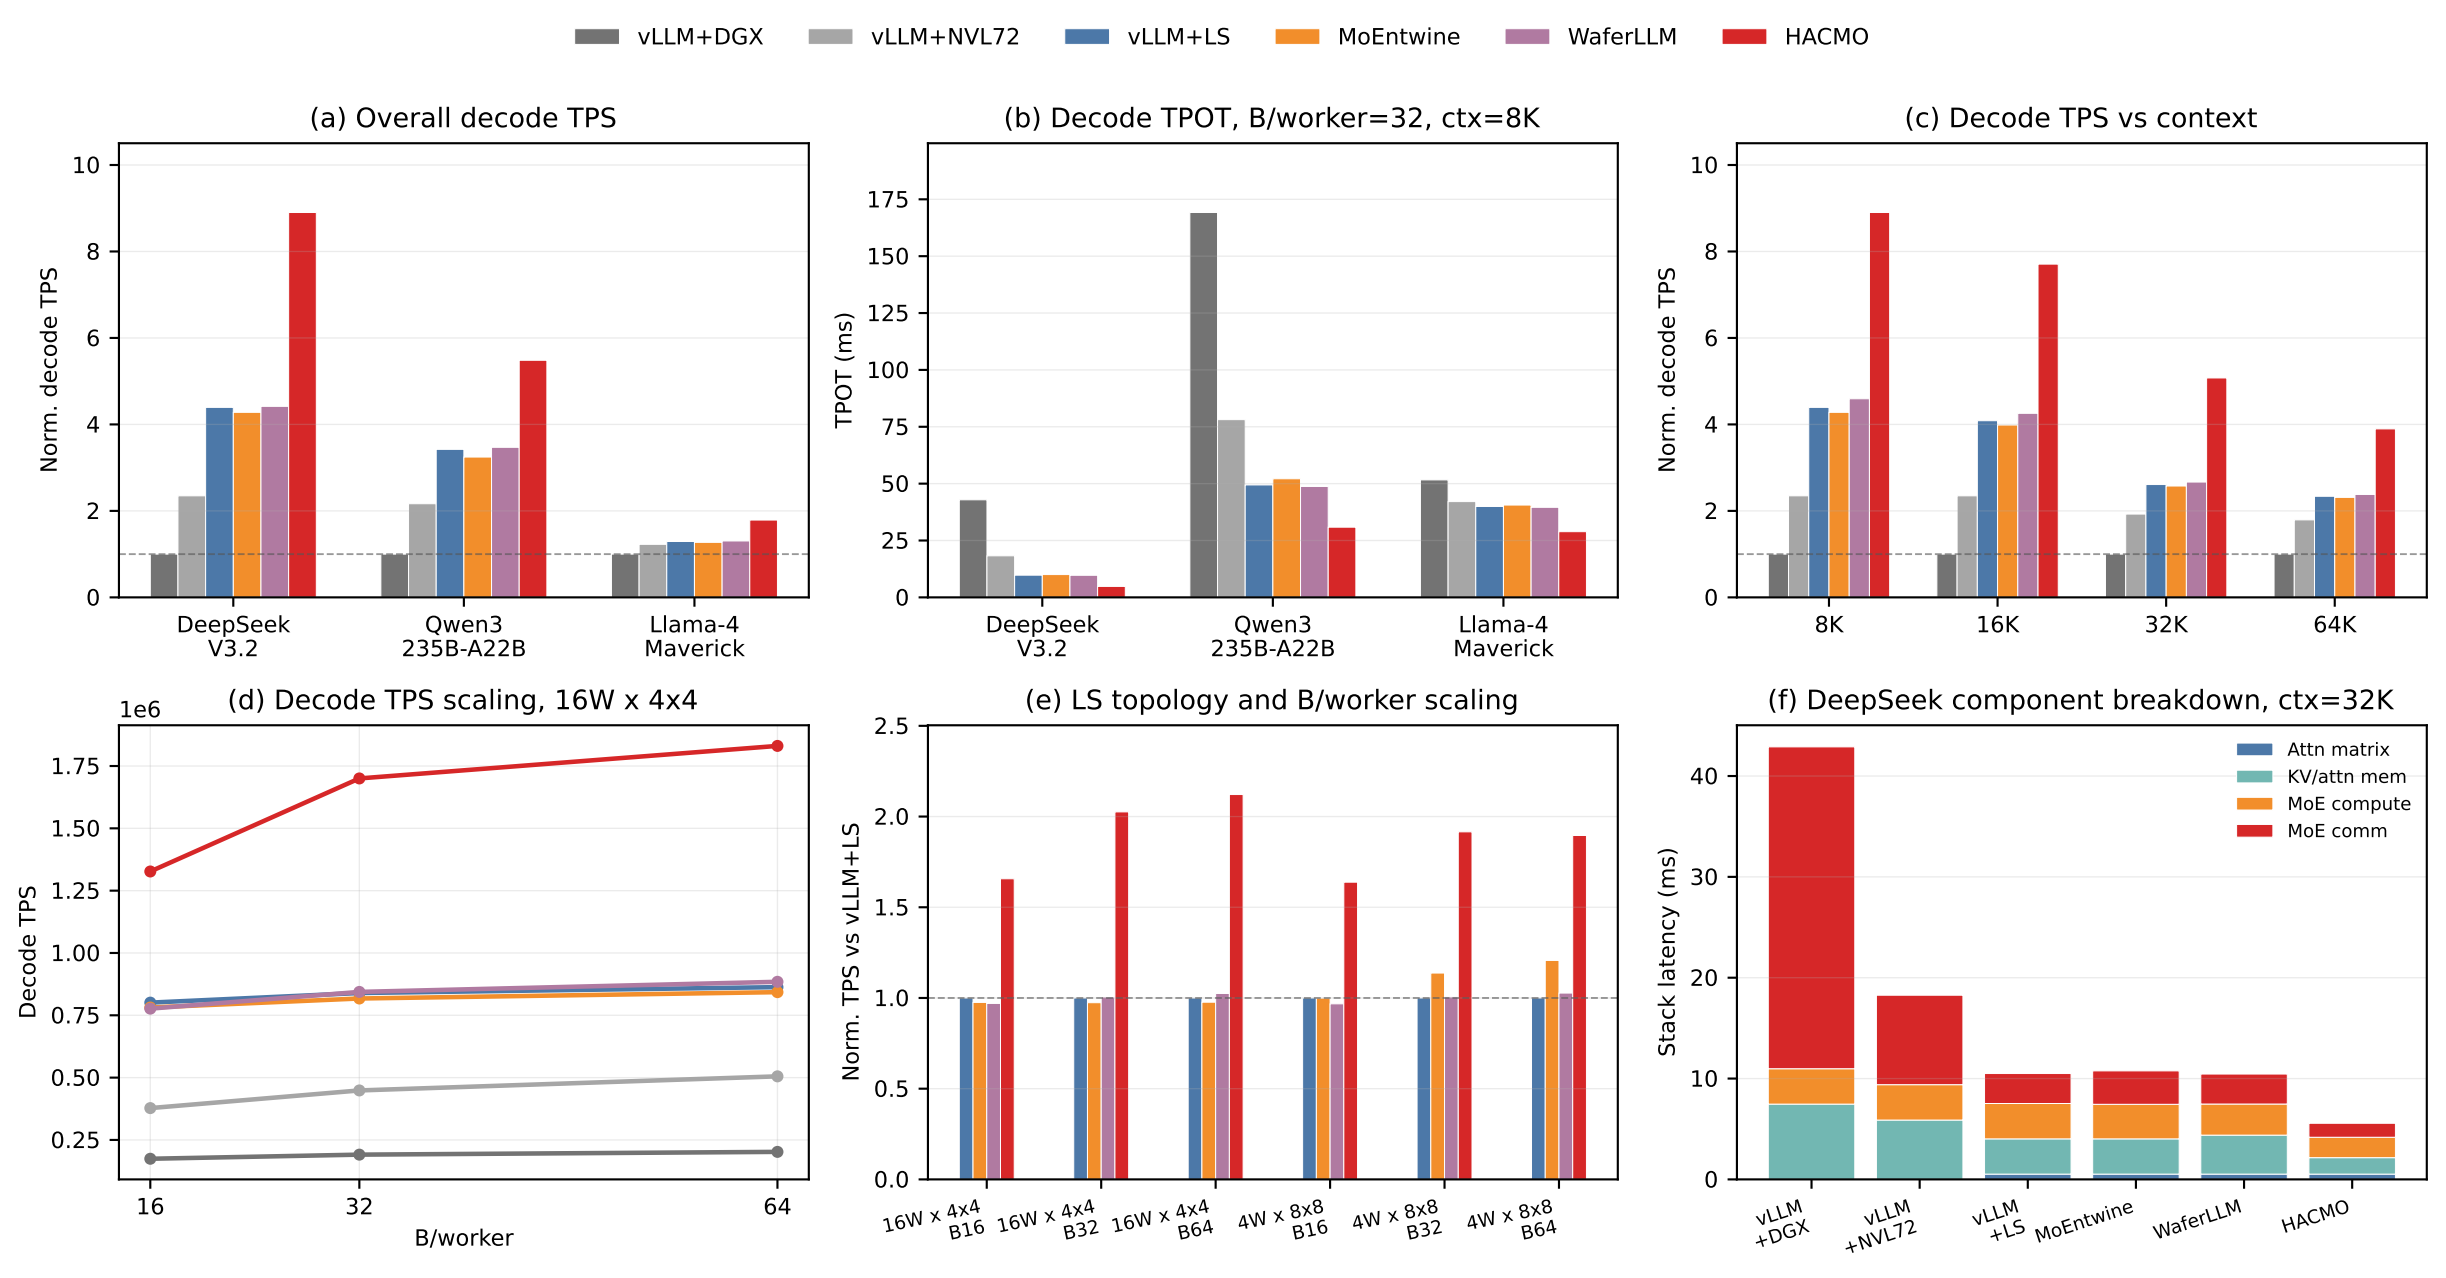
vLLM+DGX
vLLM+NVL72
vLLM+LS
MoEntwine
WaferLLM
HACMO
(a) Overall decode TPS
Norm. decode TPS
DeepSeek V3.2
Qwen3 235B-A22B
Llama-4 Maverick
(b) Decode TPOT, B/worker=32, ctx=8K
TPOT (ms)
(c) Decode TPS vs context
8K
16K
32K
64K
(d) Decode TPS scaling, 16W x 4x4
1e6
Decode TPS
B/worker
(e) LS topology and B/worker scaling
Norm. TPS vs vLLM+LS
16W x 4x4 B16
16W x 4x4 B32
16W x 4x4 B64
4W x 8x8 B16
4W x 8x8 B32
4W x 8x8 B64
(f) DeepSeek component breakdown, ctx=32K
Stack latency (ms)
Attn matrix
KV/attn mem
MoE compute
MoE comm
vLLM +DGX
vLLM +NVL72
vLLM +LS
MoEntwine
WaferLLM
HACMO